

```
In [1]: # Importing libraries

from __future__ import print_function
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn import tree
import warnings
warnings.filterwarnings('ignore')
```

```
In [12]: PATH = '/content/Crop_recommendation.csv'
df = pd.read_csv(PATH)
```

```
In [13]: df.head()
```

```
Out[13]:
```

| | N | P | K | temperature | humidity | ph | rainfall | label |
|---|----|----|----|-------------|-----------|----------|------------|-------|
| 0 | 90 | 42 | 43 | 20.879744 | 82.002744 | 6.502985 | 202.935536 | rice |
| 1 | 85 | 58 | 41 | 21.770462 | 80.319644 | 7.038096 | 226.655537 | rice |
| 2 | 60 | 55 | 44 | 23.004459 | 82.320763 | 7.840207 | 263.964248 | rice |
| 3 | 74 | 35 | 40 | 26.491096 | 80.158363 | 6.980401 | 242.864034 | rice |
| 4 | 78 | 42 | 42 | 20.130175 | 81.604873 | 7.628473 | 262.717340 | rice |

```
In [14]: df.tail()
```

```
Out[14]:
```

| | N | P | K | temperature | humidity | ph | rainfall | label |
|------|-----|----|----|-------------|-----------|----------|------------|--------|
| 2195 | 107 | 34 | 32 | 26.774637 | 66.413269 | 6.780064 | 177.774507 | coffee |
| 2196 | 99 | 15 | 27 | 27.417112 | 56.636362 | 6.086922 | 127.924610 | coffee |
| 2197 | 118 | 33 | 30 | 24.131797 | 67.225123 | 6.362608 | 173.322839 | coffee |
| 2198 | 117 | 32 | 34 | 26.272418 | 52.127394 | 6.758793 | 127.175293 | coffee |
| 2199 | 104 | 18 | 30 | 23.603016 | 60.396475 | 6.779833 | 140.937041 | coffee |

```
In [15]: df.size
```

```
Out[15]: 17600
```

```
In [16]: df.shape
```

```
Out[16]: (2200, 8)
```

```
In [17]: df.columns
```

```
Out[17]: Index(['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall', 'label'], dtype='object')
```

```
In [18]: df['label'].unique()
```

```
Out[18]: array(['rice', 'maize', 'chickpea', 'kidneybeans', 'pigeonpeas',  
              'mothbeans', 'mungbean', 'blackgram', 'lentil', 'pomegranate',  
              'banana', 'mango', 'grapes', 'watermelon', 'muskmelon', 'apple',  
              'orange', 'papaya', 'coconut', 'cotton', 'jute', 'coffee'],  
              dtype=object)
```

```
In [19]: df.dtypes
```

```
In [19]: df.dtypes
```


```
Out[19]: N          int64
P          int64
K          int64
temperature  float64
humidity     float64
ph           float64
rainfall     float64
label        object
dtype: object
```

```
In [20]: df['label'].value_counts()
```

```
Out[20]: rice          100
maize          100
jute           100
cotton         100
coconut        100
papaya         100
orange         100
apple          100
muskmelon      100
watermelon     100
grapes         100
mango          100
banana         100
pomegranate    100
lentil         100
blackgram      100
mungbean       100
mothbeans      100
pigeonpeas     100
kidneybeans     100
chickpea       100
coffee         100
Name: label, dtype: int64
```

```
In [21]: sns.heatmap(df.corr(),annot=True)
```

Out[21]:



Seperating features and target label

```
In [22]: features = df[['N', 'P','K','temperature', 'humidity', 'ph', 'rainfall']]
target = df['label']
labels = df['label']
```

```
In [23]: # Initializing empty lists to append all model's name and corresponding name
acc = []
model = []
```

```
In [24]: from sklearn.model_selection import train_test_split
Xtrain, Xtest, Ytrain, Ytest = train_test_split(features,target,test_size = 0.2,random_state =2)
```

Decision Tree

In [25]:

```
from sklearn.tree import DecisionTreeClassifier

DecisionTree = DecisionTreeClassifier(criterion="entropy",random_state=2,max_depth=5)

DecisionTree.fit(Xtrain,Ytrain)

predicted_values = DecisionTree.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Decision Tree')
print("DecisionTrees's Accuracy is: ", x*100)

print(classification_report(Ytest,predicted_values))
```

DecisionTrees's Accuracy is: 90.0

| | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| apple | 1.00 | 1.00 | 1.00 | 13 |
| banana | 1.00 | 1.00 | 1.00 | 17 |
| blackgram | 0.59 | 1.00 | 0.74 | 16 |
| chickpea | 1.00 | 1.00 | 1.00 | 21 |
| coconut | 0.91 | 1.00 | 0.95 | 21 |
| coffee | 1.00 | 1.00 | 1.00 | 22 |
| cotton | 1.00 | 1.00 | 1.00 | 20 |
| grapes | 1.00 | 1.00 | 1.00 | 18 |
| jute | 0.74 | 0.93 | 0.83 | 28 |
| kidneybeans | 0.00 | 0.00 | 0.00 | 14 |
| lentil | 0.68 | 1.00 | 0.81 | 23 |
| maize | 1.00 | 1.00 | 1.00 | 21 |
| mango | 1.00 | 1.00 | 1.00 | 26 |
| mothbeans | 0.00 | 0.00 | 0.00 | 19 |
| mungbean | 1.00 | 1.00 | 1.00 | 24 |
| muskmelon | 1.00 | 1.00 | 1.00 | 23 |
| orange | 1.00 | 1.00 | 1.00 | 29 |
| papaya | 1.00 | 0.84 | 0.91 | 19 |
| pigeonpeas | 0.62 | 1.00 | 0.77 | 18 |
| pomegranate | 1.00 | 1.00 | 1.00 | 17 |
| rice | 1.00 | 0.62 | 0.77 | 16 |
| watermelon | 1.00 | 1.00 | 1.00 | 15 |

```
In [26]: from sklearn.model_selection import cross_val_score
```

```
In [27]: # Cross validation score (Decision Tree)
score = cross_val_score(DecisionTree, features, target, cv=5)
```

```
In [28]: score
```

```
Out[28]: array([0.93636364, 0.90909091, 0.91818182, 0.87045455, 0.93636364])
```

Saving trained Decision Tree model

```
In [29]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
DT_pkl_filename = 'DecisionTree.pkl'
# Open the file to save as pkl file
DT_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(DecisionTree, DT_Model_pkl)
# Close the pickle instances
DT_Model_pkl.close()
```

```
In [30]: from sklearn.naive_bayes import GaussianNB

NaiveBayes = GaussianNB()

NaiveBayes.fit(Xtrain, Ytrain)

predicted_values = NaiveBayes.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Naive Bayes')
print("Naive Bayes's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))
```

```
Naive Bayes's Accuracy is: 0.990909090909091
      precision    recall  f1-score   support
```

```
In [39]: # Cross validation score (Random Forest)
score = cross_val_score(RF, features, target, cv=5)
score
```

```
Out[39]: array([0.99772727, 0.99545455, 0.99772727, 0.99318182, 0.98863636])
```

```
In [40]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
RF_pkl_filename = 'RandomForest.pkl'
# Open the file to save as pkl file
RF_Model_pkl = open(RF_pkl_filename, 'wb')
pickle.dump(RF, RF_Model_pkl)
# Close the pickle instances
RF_Model_pkl.close()
```

```
In [41]: import xgboost as xgb
XB = xgb.XGBClassifier()
XB.fit(Xtrain, Ytrain)

predicted_values = XB.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('XGBoost')
print("XGBoost's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))
```


```
XGBoost's Accuracy is: 0.9931818181818182
```

| | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| apple | 1.00 | 1.00 | 1.00 | 13 |
| banana | 1.00 | 1.00 | 1.00 | 17 |
| blackgram | 1.00 | 1.00 | 1.00 | 16 |
| chickpea | 1.00 | 1.00 | 1.00 | 21 |
| coconut | 1.00 | 1.00 | 1.00 | 21 |
| coffee | 1.00 | 1.00 | 1.00 | 22 |
| cotton | 1.00 | 1.00 | 1.00 | 20 |
| grapes | 1.00 | 1.00 | 1.00 | 18 |
| jute | 0.96 | 0.93 | 0.95 | 28 |
| kidneybeans | 1.00 | 1.00 | 1.00 | 14 |

```
# Dump the trained Naive Bayes classifier with Pickle
XB_pkl_filename = 'XGBoost.pkl'
# Open the file to save as pkl file
XB_Model_pkl = open(XB_pkl_filename, 'wb')
pickle.dump(XB, XB_Model_pkl)
# Close the pickle instances
XB_Model_pkl.close()
```

```
In [43]: plt.figure(figsize=[10,5],dpi = 100)
plt.title('Accuracy Comparison')
plt.xlabel('Accuracy')
plt.ylabel('Algorithm')
sns.barplot(x = acc,y = model,palette='dark')
```

Out[43]:



accuracy

```
In [44]: accuracy_models = dict(zip(model, acc))
         for k, v in accuracy_models.items():
             print (k, '-->', v)
```

```
Decision Tree --> 0.9
Naive Bayes --> 0.990909090909091
SVM --> 0.10681818181818181
Logistic Regression --> 0.9522727272727273
RF --> 0.990909090909091
XGBoost --> 0.9931818181818182
```

```
In [45]: data = np.array([[104,18, 30, 23.603016, 60.3, 6.7, 140.91]])
         prediction = RF.predict(data)
         print(prediction)
```

['coffee']

```
In [46]: data = np.array([[83, 45, 60, 28, 70.3, 7.0, 150.9]])
         prediction = RF.predict(data)
         print(prediction)
```

['jute']